

PureGym Support Hub

Project overview, operating model, and cost estimate

Updated July 4, 2026

Executive Summary

PureGym Support Hub is an internal support workspace for KSA and UAE agents. It centralizes approved scripts, quick replies, branch references, admin control, AI assistance, and operational checks in one Vercel-hosted dashboard backed by Neon PostgreSQL.

Latest Delivered Updates

- AI provider is no longer tied to OpenAI. Groq is the default low-cost provider, with Gemini and OpenAI kept as optional fallbacks.
- Ctrl+K now opens the in-app command palette when the page has focus, making script search faster than navigating categories.
- Fill-before-copy variables now work across scripts that contain dates, duration, membership type, amount, or member name placeholders.
- Smoke tests were added for the dashboard shell, admin protection, script APIs, and quick-script access rules.
- Documentation and cost assumptions now match the current Vercel, Neon, Groq, and Gemini setup.

Primary Capabilities

Agent Workspace

- Script Library
- Quick Scripts
- Favorites
- Branch Directory

Admin Control

- Script/category activation
- Quick-script management
- User approval
- Role separation

AI Assistance

- Support chatbot
- Spellcheck helper
- AI Trainer references
- Memory summaries

Production Architecture

- Frontend and API routes run on Next.js 14 App Router deployed to Vercel.
- Persistent application data lives in Neon PostgreSQL through Prisma.
- Scripts edited in the admin dashboard are stored in Neon, so deployment does not overwrite live admin additions.
- Rate limiting protects login, signup, password reset, and AI chat endpoints.
- Security headers and httpOnly signed session cookies are enabled for production use.

Roles and Data Safety

- Super Admin controls user role changes and account removal.
- Admin can manage operational scripts and approve/generate passwords for new users.
- Regular users can use scripts, quick replies, favorites, calculations, and AI without changing shared content.
- Category and script deactivation hides content globally without deleting historical data.
- Legacy quick-script data is audited before cleanup, and cleanup only runs when rows are confirmed duplicates.

QA and Operational Checks

- `npm run build` validates Prisma generation, Next.js compilation, linting, type checks, and page generation.
- `npm run smoke` validates public page rendering and anonymous API protections by default.
- Authenticated smoke checks can be enabled with `SMOKE_EMAIL` and `SMOKE_PASSWORD`.
- Authenticated smoke checks can verify shared quick scripts without creating legacy per-user copies.
- `/api/health` supports external monitoring and can be protected with `HEALTH_SECRET` or `CRON_SECRET`.

Expected Monthly Cost for 25 Daily Users

The estimate below assumes 25 support users, daily script usage, moderate AI chat usage, and one developer/admin seat managing deployment. Support agents are application users, not Vercel developer seats.

Vercel hosting	\$0 on Hobby for light use; \$20/mo Pro recommended for production ownership and higher limits.
Neon PostgreSQL	\$0 Free can work for small data; Launch typical spend is about \$15/mo if Free limits are exceeded.
AI provider	Groq Llama 3.1 8B is very low cost at \$0.05/M input and \$0.08/M output tokens; Gemini API also has free usage through Google AI Studio.
Email SMTP	Usually \$0 if using an existing Gmail/SMTP account within provider limits.
Practical range	\$0-\$10/mo for light internal use, or roughly \$35-\$45/mo for a safer production setup with Vercel Pro + Neon Launch + AI buffer.

Development Roadmap

- Browser extension connected to the hub, replacing TextBlaze-style snippets with centrally managed shortcuts.
- Health endpoint monitors AI configuration, database access, cron configuration, and offer-sync freshness.
- Legacy database audit/cleanup command removes UserQuickScript rows only when they are confirmed duplicates.

Sources

Vercel pricing: <https://vercel.com/pricing>
Neon pricing: <https://neon.com/pricing>
Groq pricing: <https://groq.com/pricing>
Gemini API pricing: <https://ai.google.dev/gemini-api/docs/pricing>