

PureGym Support Hub

Project overview, operating model, and feature coverage

Updated July 4, 2026

Executive Summary

PureGym Support Hub is an internal support workspace for KSA and UAE agents. It centralizes approved scripts, quick replies, branch references, admin control, AI assistance, and operational checks in one Vercel-hosted dashboard backed by Neon PostgreSQL.

Latest Delivered Updates

- AI provider is no longer tied to OpenAI. Groq is the default provider, with Gemini and OpenAI kept as optional fallbacks.
- Ctrl+K now opens the in-app command palette when the page has focus, making script search faster than navigating categories.
- Fill-before-copy variables now work across scripts that contain dates, duration, membership type, amount, or member name placeholders.
- Smoke tests were added for the dashboard shell, admin protection, script APIs, and quick-script access rules.
- Documentation now matches the current Vercel, Neon, Groq, and Gemini setup.

Primary Capabilities

Agent Workspace

- Script Library
- Quick Scripts
- Favorites
- Branch Directory

Admin Control

- Script/category activation
- Quick-script management
- Shared content controls
- Content status management

AI Assistance

- Support chatbot
- Spellcheck helper
- AI Trainer references
- Memory summaries

Production Architecture

- Frontend and API routes run on Next.js 14 App Router deployed to Vercel.
- Persistent application data is stored in Neon PostgreSQL through Prisma.
- Scripts edited in the admin dashboard are stored in Neon, so deployment does not overwrite current admin additions.
- Rate limiting protects login, signup, password reset, and AI chat endpoints.
- Security headers and httpOnly signed session cookies are enabled for production use.

Roles and Data Safety

- Admin can manage operational scripts, categories, Quick Scripts, activation status, and shared support content.
- User account creation, removal, and role changes are left to the technical team rather than handled in this presentation scope.
- Regular users can use scripts, quick replies, favorites, calculations, and AI without changing shared content.
- Category and script deactivation hides content globally without deleting historical data.
- Legacy quick-script data is audited before cleanup, and cleanup only runs when rows are confirmed duplicates.

QA and Operational Checks

- `npm run build` validates Prisma generation, Next.js compilation, linting, type checks, and page generation.
- `npm run smoke` validates public page rendering and anonymous API protections by default.
- Authenticated smoke checks can be enabled with `SMOKE_EMAIL` and `SMOKE_PASSWORD`.
- Authenticated smoke checks can verify shared quick scripts without creating legacy per-user copies.
- `/api/health` supports external monitoring for database, AI configuration, and system readiness.

Workspace Features

- Script Library with category filters, active/inactive visibility for admins, favorites, and one-click copy behavior.
- Quick Scripts with shared admin-managed content, drag-and-drop ordering, personal favorites, and male/female wording toggles.
- Ctrl+K command palette for fast script lookup and keyboard-first copying.
- Fill-before-copy prompts for dates, duration, membership type, amount, and member names.
- Branch Directory with searchable KSA/UAE branch references.

Assistant and Admin Features

- AI chatbot uses official knowledge, scripts, trainer notes, conversation memory, and country/language signals.
- AI Trainer accepts text, links, and image-based references without forcing agents to tag every item manually.
- Spellcheck helper improves selected scripts while keeping the agent in control.
- Admin Editor manages scripts, categories, activation status, account approvals, and generated passwords.
- Calculation Tool includes membership calculations and Hijri/Gregorian conversion in both directions.

Reliability Features

- Shared Quick Scripts now read from the master Script table, preventing old per-user copies from returning.
- Legacy UserQuickScript audit refuses cleanup if any unique user content is detected.
- Smoke tests cover public pages, route protection, health checks, and authenticated flows when credentials are provided.
- Health checks expose database, AI configuration, and system readiness for external monitors.

Monthly Cost Model

The hub is designed for low operating cost because the main workload is an internal Next.js dashboard, a small PostgreSQL database, and AI calls that run only when agents use the assistant or helper actions. The numbers below are planning estimates for about 25 application users, not final invoices.

Cost Details

- Vercel hosting: \$0/mo on Hobby for light internal use, or \$20/mo on Pro with 1 deploying team seat and \$20 monthly usage credit. This pays for the Next.js app, API routes, CDN, builds, deployments, and production controls.
- Neon PostgreSQL: \$0/mo can work for a small database; Launch is a typical \$15/mo for intermittent load and 1 GB. This pays for persistent scripts, categories, favorites, account requests, AI references, and memory summaries.
- Groq AI: llama-3.1-8b-instant is \$0.05 per 1M input tokens and \$0.08 per 1M output tokens. AI cost appears only when chat, spellcheck, trainer extraction, or summary actions are used.
- Gemini optional fallback: selected Gemini API models include a free tier, so it can be used as a no-cost path for light AI usage within rate limits.
- Email, domain, and basic monitoring: \$0/mo when existing SMTP, the Vercel domain, and built-in checks are enough. Custom domains, paid email, or external monitoring depend on the provider selected.

Practical Budget

- Lean internal setup: about \$0-\$10/mo if Vercel Hobby, Neon Free, and low AI volume stay within plan limits.
- Recommended production baseline: about \$35-\$45/mo using Vercel Pro at \$20/mo, Neon Launch around \$15/mo, and a small AI buffer.
- 25-user AI example: 25 users x 20 AI messages/day x 30 days = 15,000 AI calls/month. At 1,500 input tokens and 700 output tokens per call on Groq, the AI token cost is about \$2/mo.
- A heavier AI workload can be budgeted at \$5-\$10/mo first, then adjusted after real usage appears in provider dashboards.

Cost Notes

- The 25 support users are application users; they are not Vercel paid team seats.
- Only people who deploy or manage the Vercel project need paid Vercel team access on Pro.
- Official cost references: Vercel plan pages, Neon plans, Groq model cost page, and Gemini API cost page.